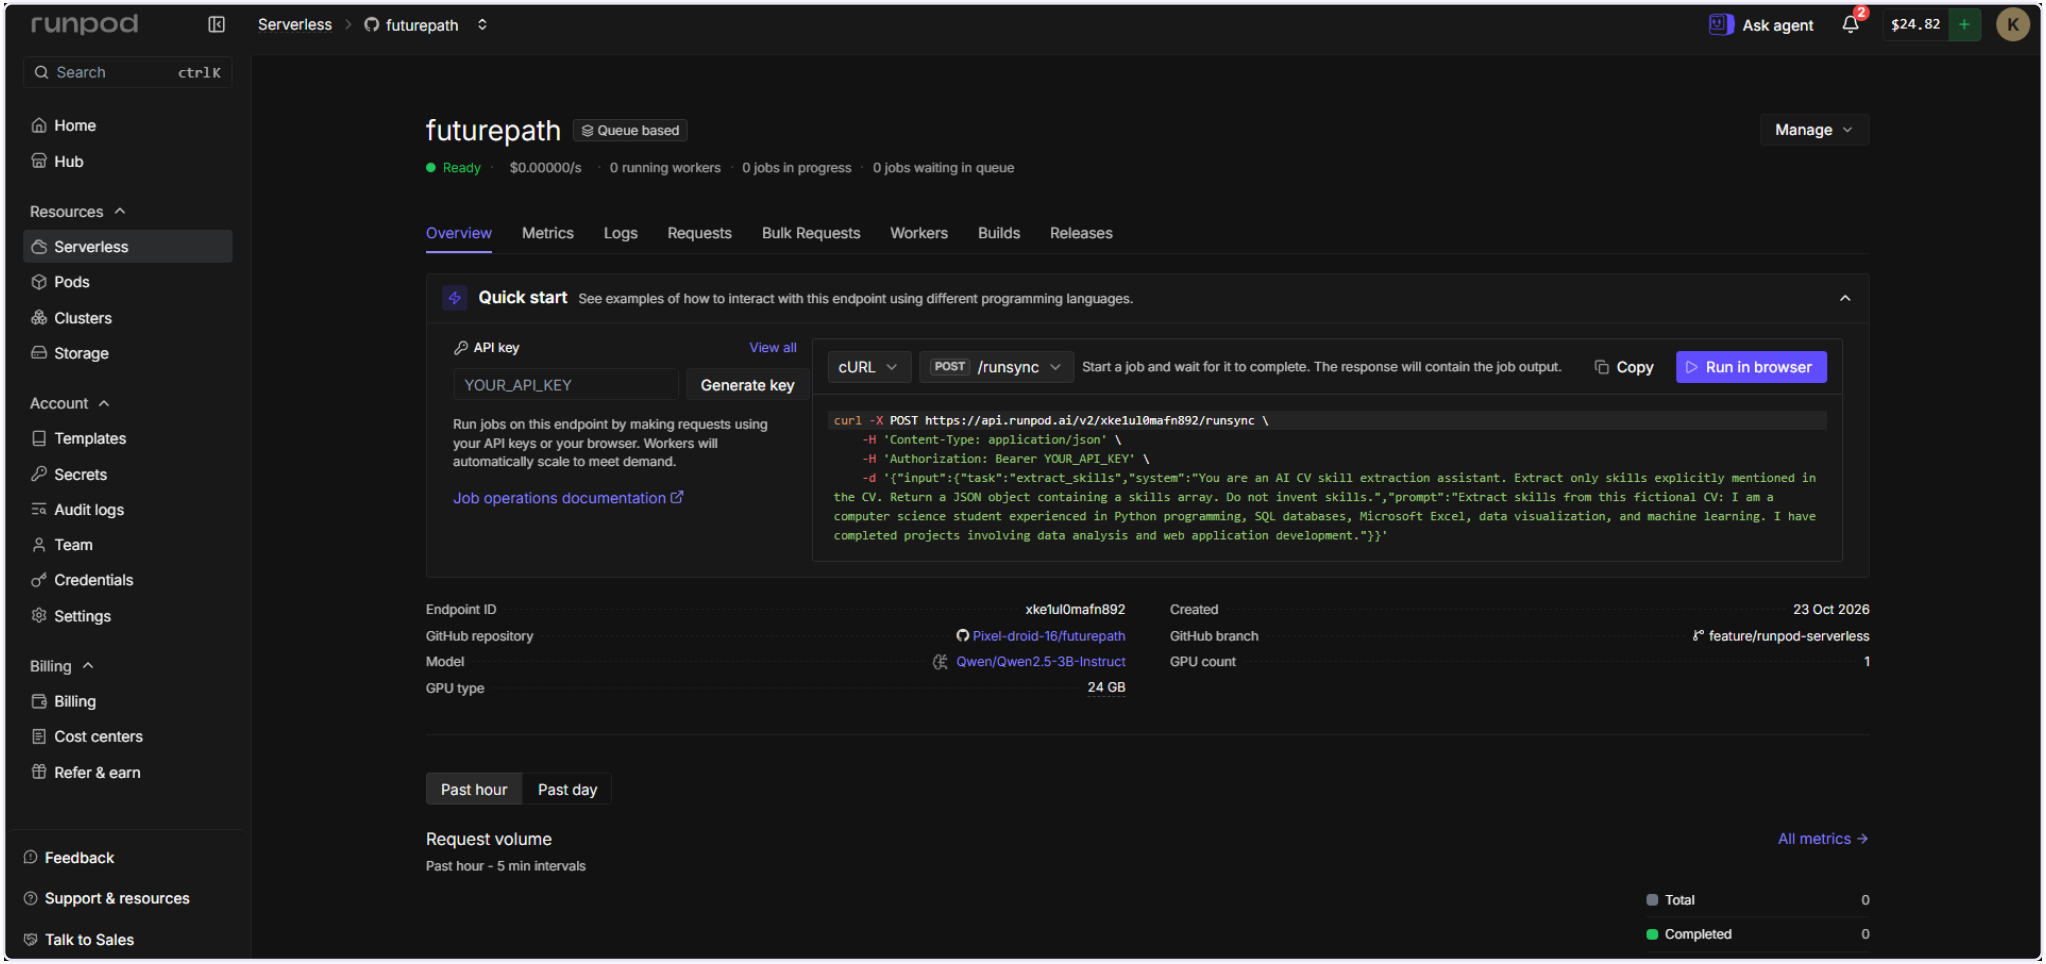
RunPod Serverless endpoint and model

EVIDENCE 01 / CONFIGURED ENDPOINT

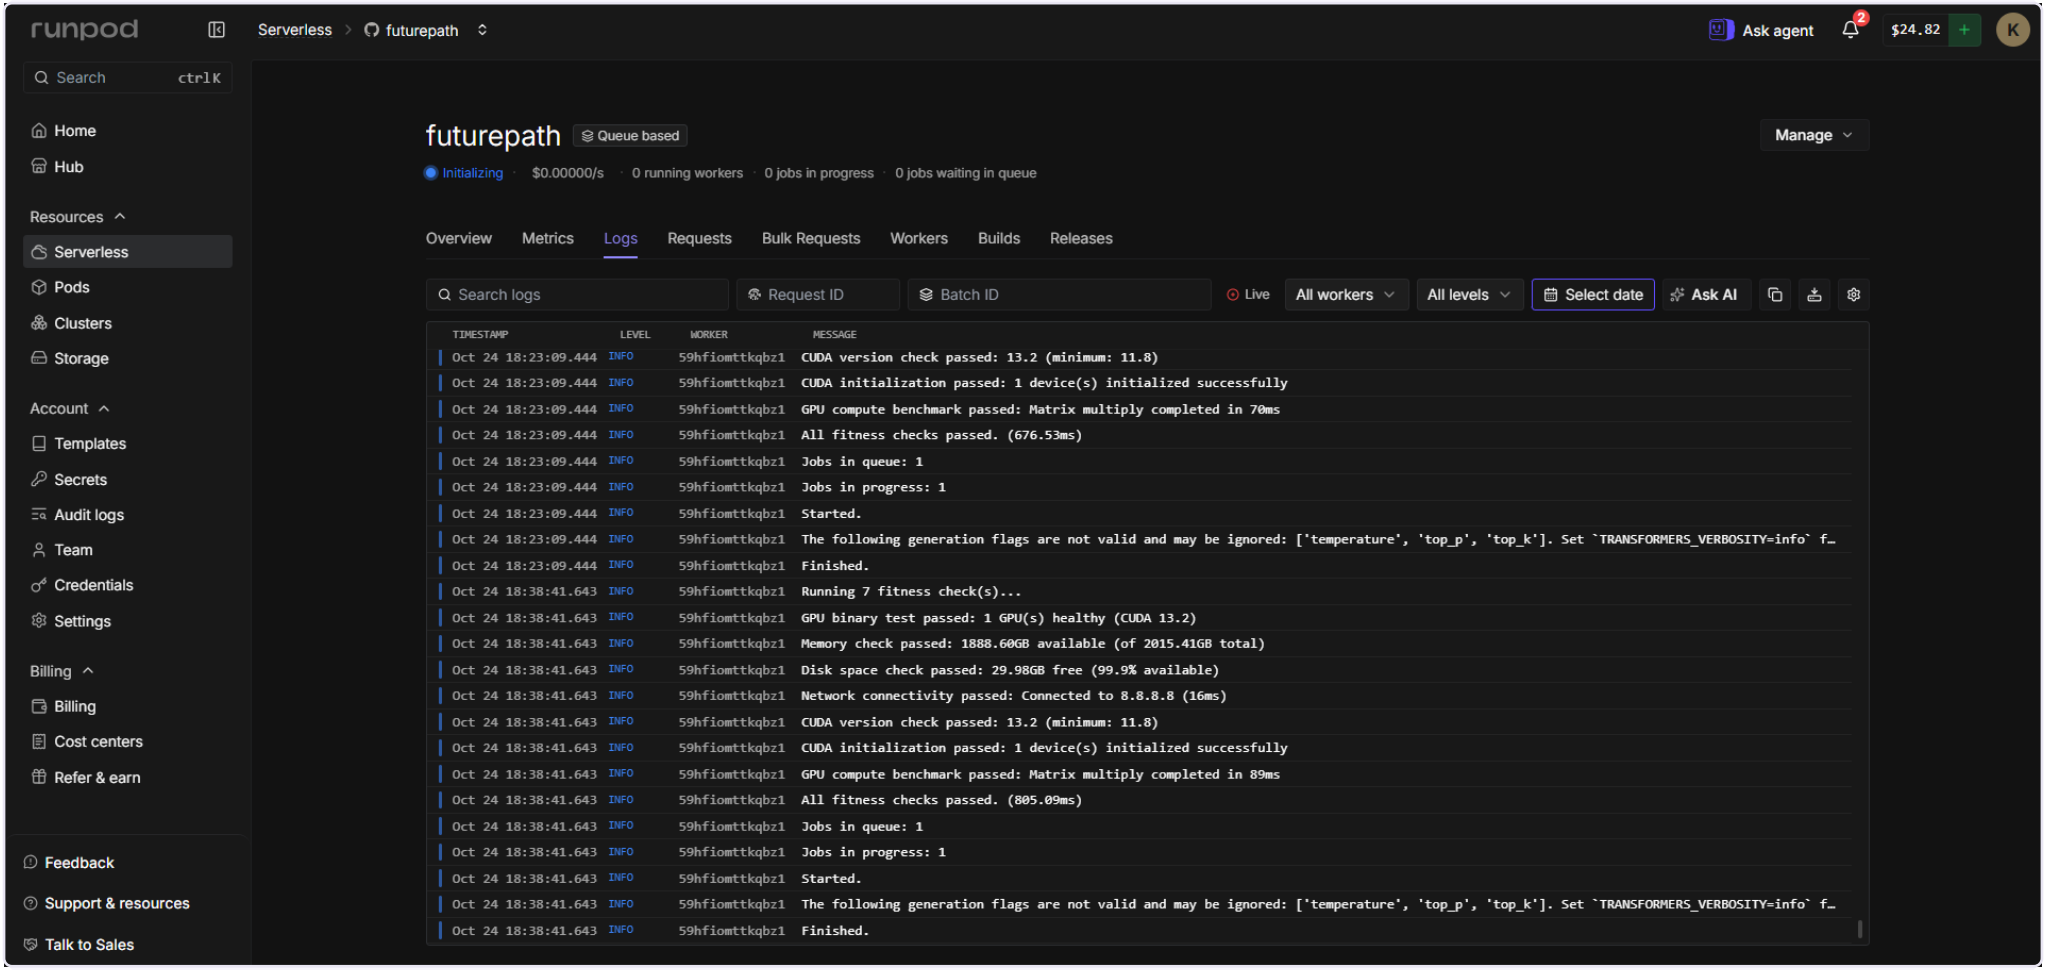
Ready endpoint linked to FuturePath GitHub and the Qwen2.5-3B-Instruct model.



RunPod worker and CUDA health logs

EVIDENCE 02 / GPU EXECUTION LOGS

GPU initialization and health checks passed; visible jobs started and finished.

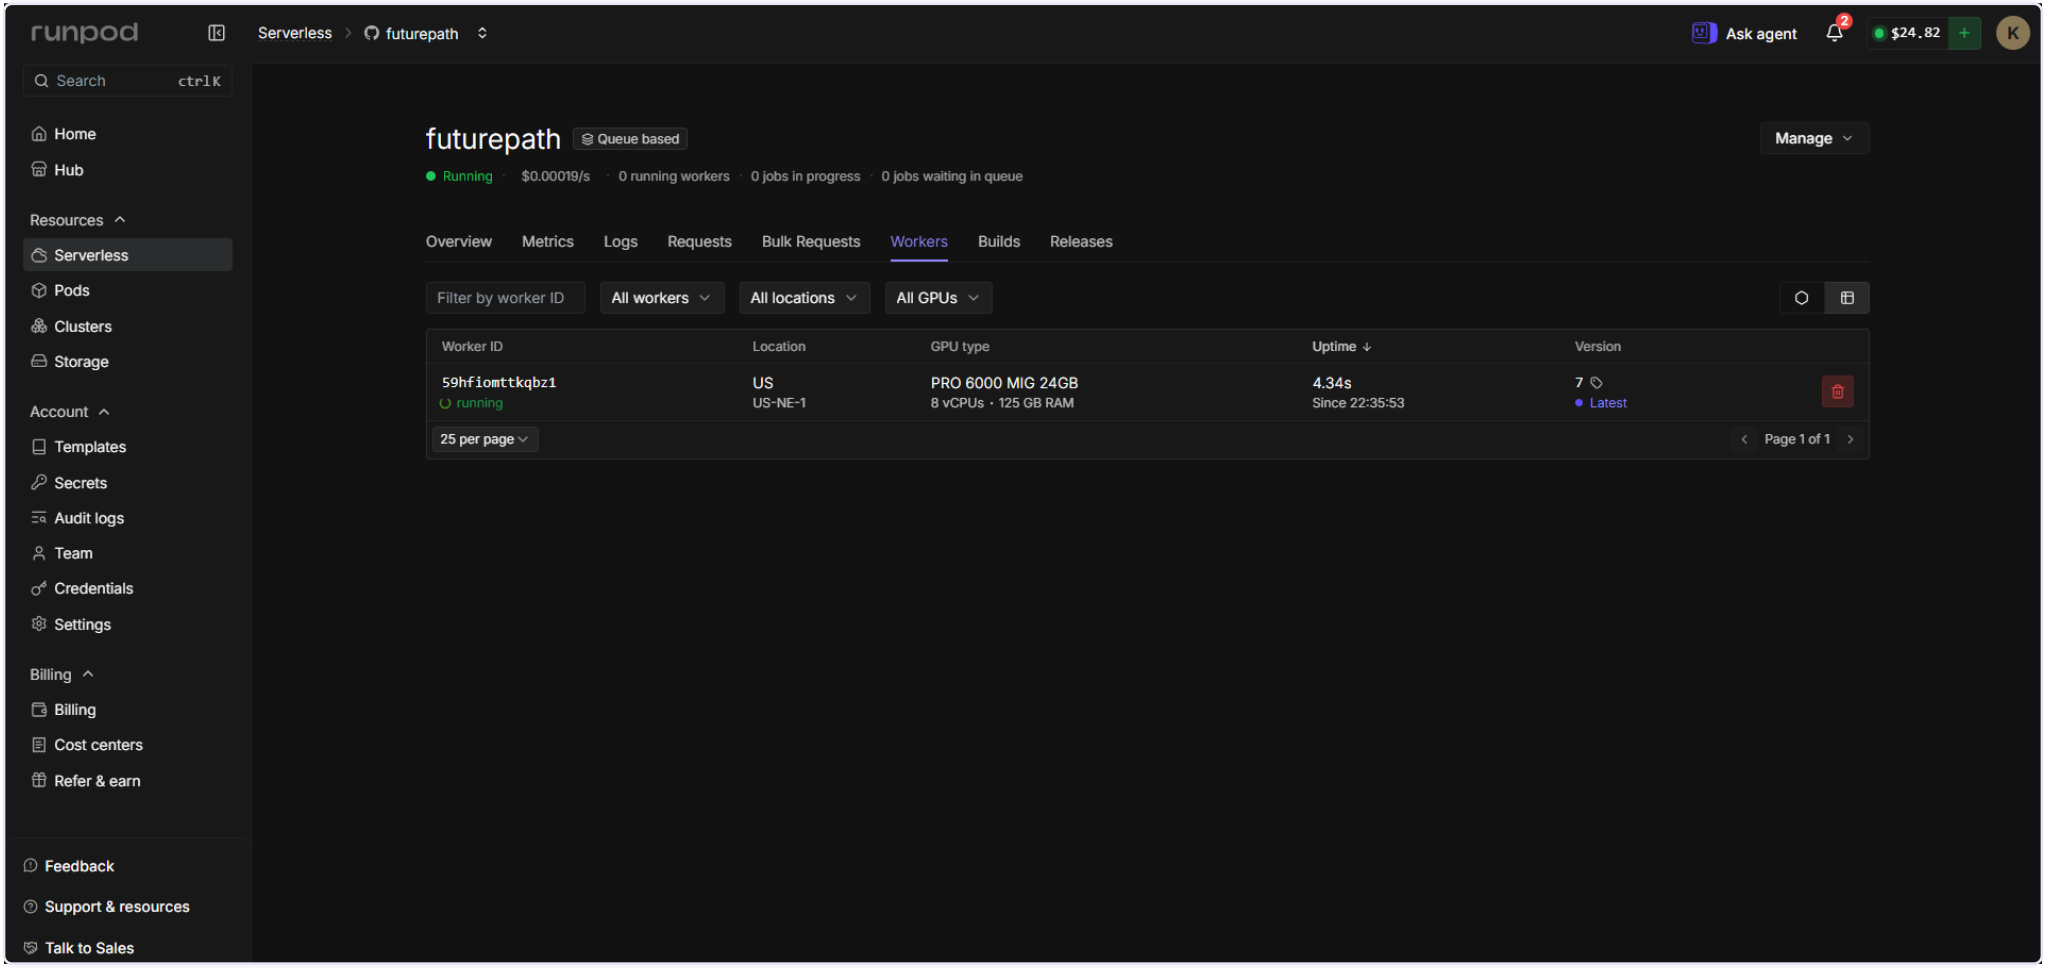


Logs show completed jobs; CV extraction output has not been independently verified.

RunPod GPU worker configuration

EVIDENCE 03 / ACTIVE GPU WORKER

Running PRO 6000 MIG 24GB worker in US-NE-1, provisioned for inference.



Source: genuine RunPod screenshots provided by the FuturePath team, 24 Oct 2026.